

Fusion of Region-Based Representations for Gender Identification

Si ying Diana Hu, Brendan Jou, Aaron Jaech, Marios Savvides



Goals & Motivations

- Create a classification framework robust to uncontrolled, real world images while avoiding over fitting
- Explore localized classification on facial discriminative regions to increase accuracy rate that can overcome:
 - Pose
 - Alignment
 - Facial expressions
 - Occlusions

Dataset

- Database comprises 26,766 (13383 male/female) unique faces collected from Flickr by the CMU Biometrics Center
- Workers on Amazon Mechanical Turk landmarked key points on the face and provided gender ground truth
- High variation in pose, illumination, expression, resolution
- Preprocessing performed to normalize rotation and scale variation




Fig: Seven facial landmark points




Fig: Point cloud distribution




Fig: Sample images from database

Region Selection & Localization


- Local regions improves alignment in each region
- These regions are non-overlapping, so they can be manipulated individually and treated as orthogonal
- Regions selected based on:
 - Discriminant ability of SVM weight vector
 - Fisher Discriminant Analysis (LDA)
 - Perceptual and psychological studies by Brown et al.
- Bounding boxes are heuristically determined facial ratios



Fig: Fisher , SVM, weight vector (Top to bottom)



Region Localization



Region Crops

Fig: Regions of Interest

Implementation

- Gender features are well represented through texture
 - MR8 filter bank robust to rotation and scaling
- A separate linear SVM is trained on each of the regions
- The SVM classifications are fused with
 - Majority voting scheme
 - Confidence of SVM margin fit to a logistic function
 - Naïve Bayes classifier trained on margin distance
 - Logistic regression trained on margin distance
 - Taking the raw sum of distances to the margin

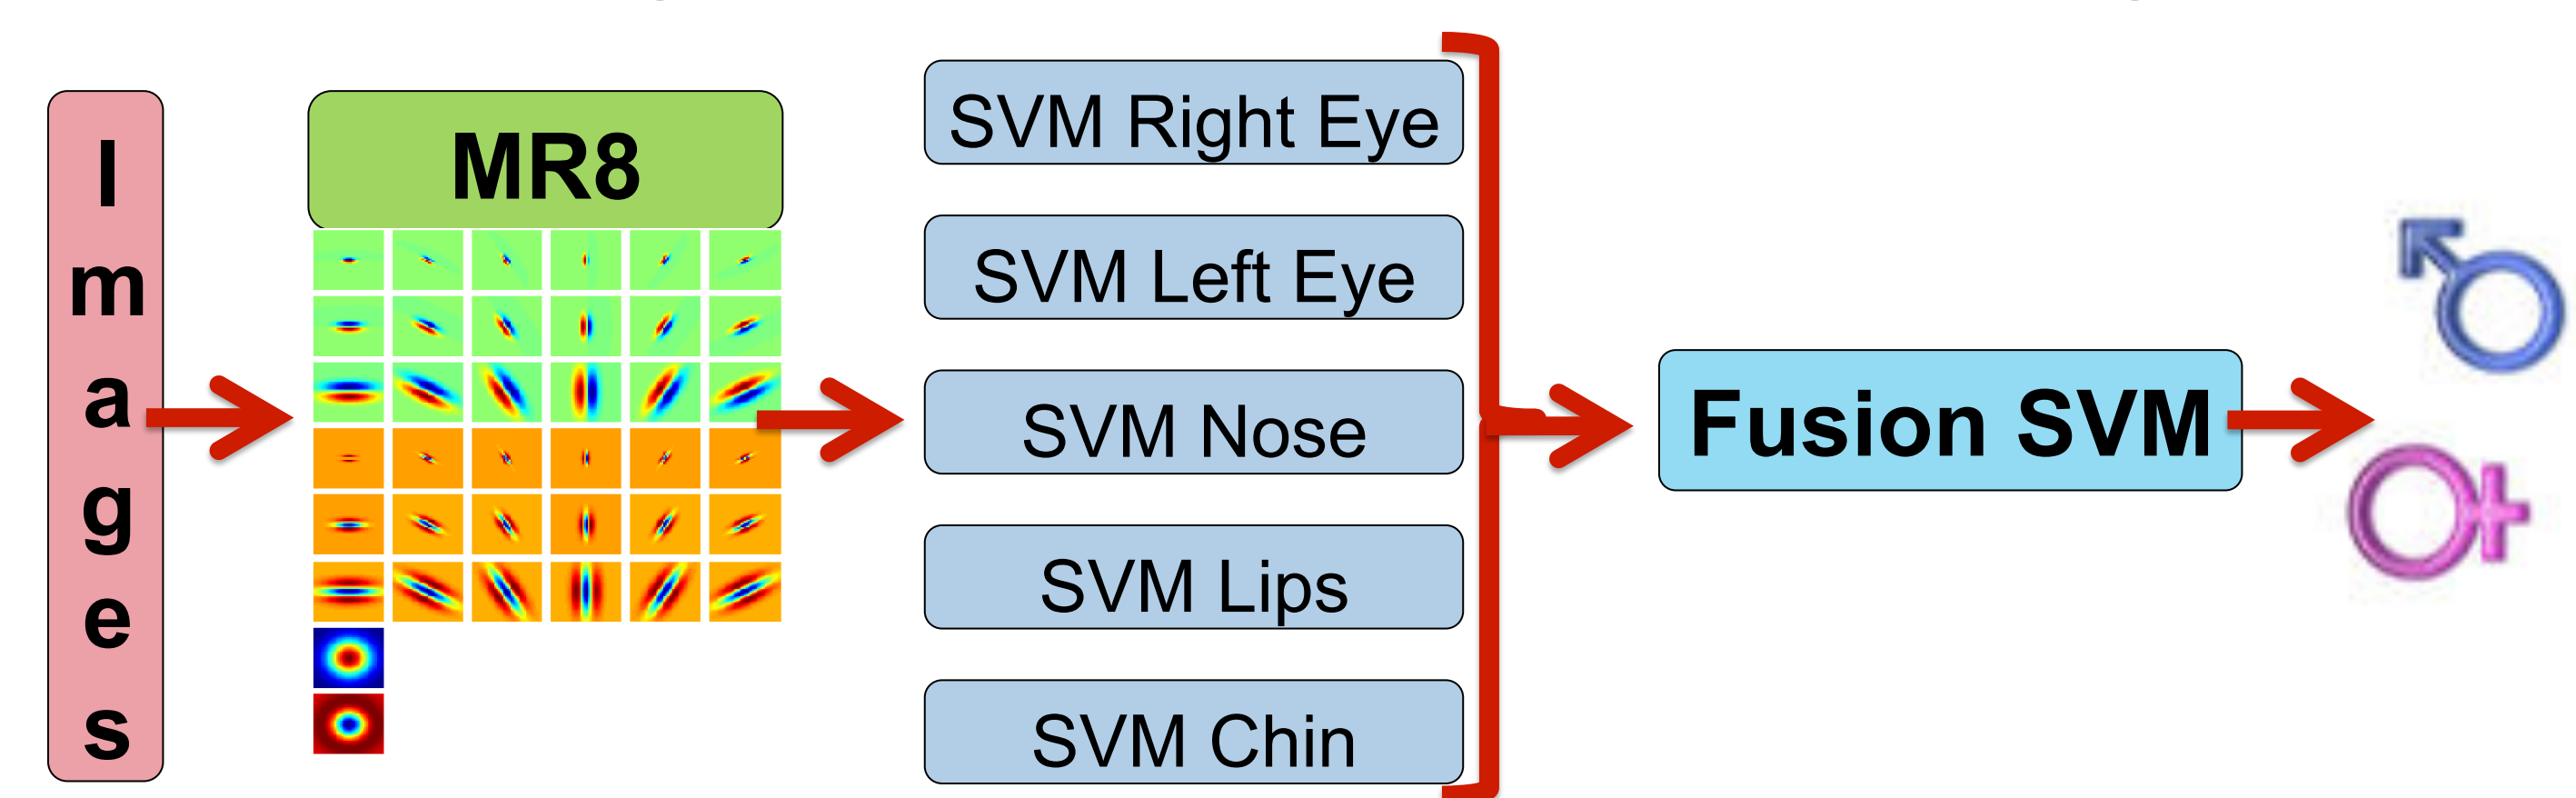


Fig: Algorithm implementation

Results & Conclusion

- Results reported on 5-fold cross validation for each dataset
- Despite the significant loss of information for each ROI, each still has significant discriminant power
- The SVM distance margins in each regions provides independent prediction of gender
 - Results in a 5 dimensional feature space
 - Likely to better separate due to region localization

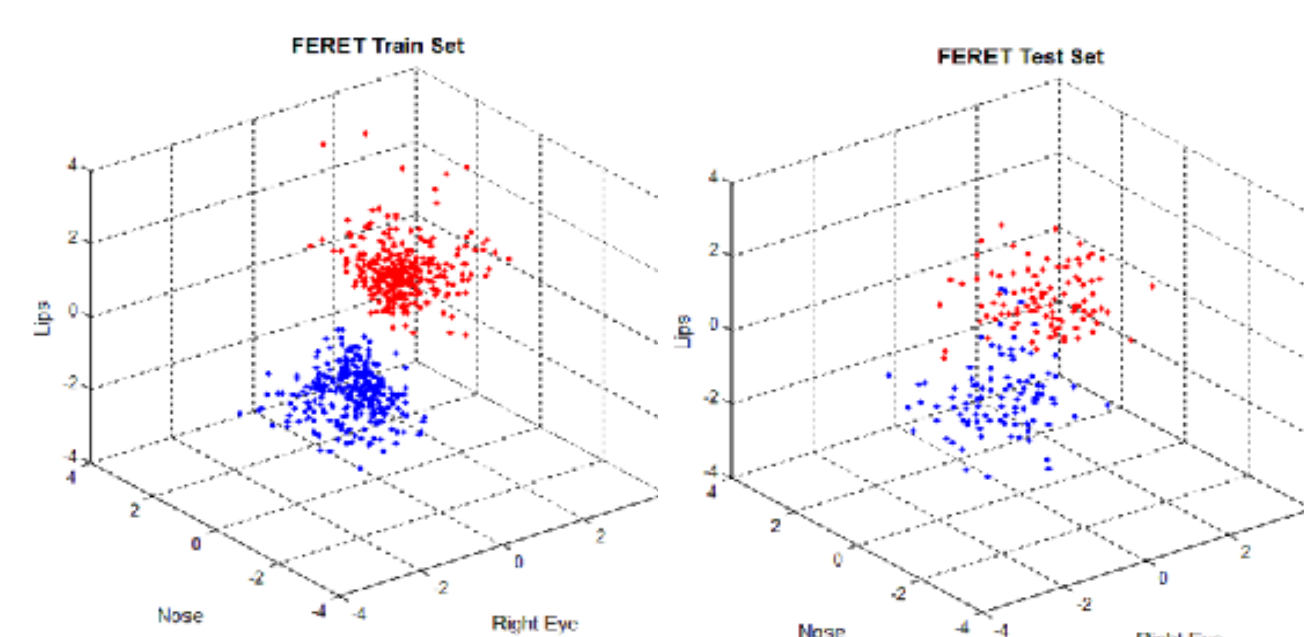


Fig: Visualization FERET score distribution

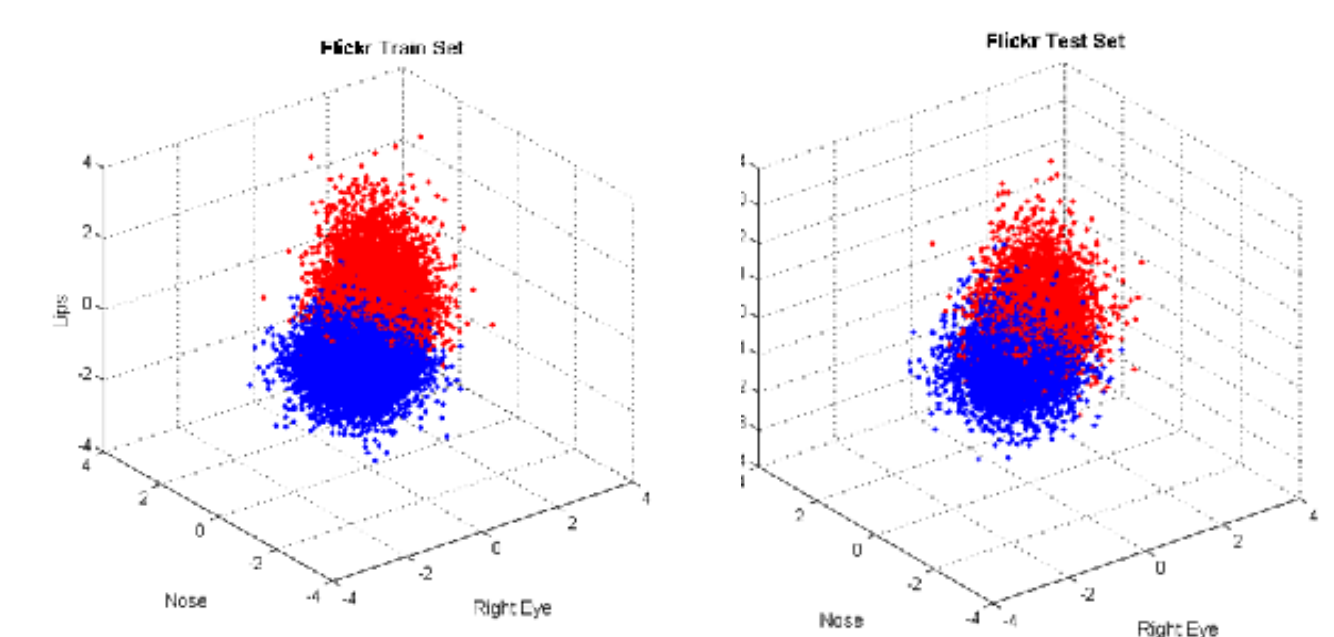


Fig: Visualization Flickr score distribution

Region	Pixel Feature		MR8 Feature		Baseline
	FERET	Flickr	FERET	Flickr	
Full face	86.6%	78.5%	90.2%	85.4%	
Right eye	85.8%	73.2%	73.8%	81.1%	
Left eye	84.2%	69.0%	75.9%	81.8%	
Nose	79.9%	66.1%	72.1%	67.7%	
Lips	83.4%	73.2%	72.7%	78.9%	
Chin	73.7%	65.7%	68.4%	76.6%	

Method	FERET	Flickr	Fusion
No Regions	90.1%	85.4%	
Majority Vote	90.9%	87.5%	
Naive Bayes	91.6%	90.1%	
Logistic Regression	92.8%	89.8%	
SVM Distances	91.9%	90.1%	

